

# **Introduction to Artificial Intelligence**

**Lecture: Adversarial Search**

# Outline

---

- Games
- Optimal Decisions in Games
- $\alpha$ - $\beta$  Pruning
- Imperfect, Real-time Decisions
- Stochastic Games

# Multiagent environments

- Each agent needs to consider the actions of other agents and how they affect its own welfare.
- The unpredictability of other agents introduce contingencies into the agent's problem-solving process
- Game theory views any multiagent environment as a game.
  - The impact of each agent on the others is “significant,” regardless of whether the agents are cooperative or competitive.
- Types of games:
  - Perfect information vs Imperfect information
  - Deterministic vs Chance

*A sequential game has perfect information if each player, when making any decision, is perfectly informed of all the events that have previously occurred, including the "initialization event" of the game.*

# Types of game

	Deterministic	Chance
Perfect information	Chess, Checkers, Go, Othello	Backgammon, Monopoly
Imperfect information		Bridge, poker, scrabble nuclear war



# Adversarial search

- Adversarial search (known as games) covers competitive environments in which the agents' goals are in conflict.
- Zero-sum games of perfect information
  - Deterministic, fully observable environments, turn-taking, two-player
  - The utility values at the end are always equal and opposite.

# Primary assumptions

- Two players only, called MAX and MIN.
  - MAX moves first, and then they take turns moving until the game ends
  - Winner gets reward, loser gets penalty.
- Both players have complete knowledge of the game's state
- No element of chance
- Zero-sum games
  - The total payoff to all players is the same for every game instance.
- Rational players
  - Each player always tries to maximize his/her utility

# Games as search


- $S_0$  – Initial state: How the game is set up at the start 0.
- $PLAYER(s)$ : Which player has the move in a state, MAX/MIN?
- $ACTIONS(s)$  – Successor function: A list of (move, state) pairs specifying legal moves.
- $RESULT(s, a)$  – Transition model: Result of move  $a$  on state  $s$
- $TERMINAL$  –  $TEST(s)$ : Is the game finished?
- $UTILITY(s,p)$  – Utility function: A numerical value of a terminal state  $s$  for a player  $p$

# Games vs. Search problems

---

- Complexity: games are too hard to be solved
- Time limits: make some decision even when calculating the optimal decision is infeasible
- Efficiency: penalize inefficiency severely
  - Several interesting ideas on how to make the best possible use of time are spawn.

# Search Tree of Tic-Tac-Toe



# Optimal decision in games


- Normal search problem
  - Optimal solution is a sequence of action leading to a goal state.
- Games
  - A search path that guarantee win for a player
  - The optimal strategy can be determined from the minimax value of each node
- $\text{MINIMAX}(s) =$

$$\left\{ \begin{array}{ll} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \end{array} \right.$$

# Example

MAX

MIN



Utility values for MAX

# The minimax algorithm

- Compute the minimax decision from the current state
- Use a simple recursive computation of the minimax values of each successor state
  - The recursion proceeds all the way down to the leaves of the tree, and then the minimax values are backed up through the tree as the recursion unwinds.

# The minimax algorithm

```
function MINIMAX-DECISION(state) returns an action  
    return argmaxa ∈ ACTIONS(s) MIN-VALUE(RESULT(state, a))
```


```
function MAX-VALUE(state) returns a utility value  
    if TERMINAL-TEST(state) then return UTILITY(state)  
    v ← -∞  
    for each a in ACTIONS(state) do  
        v ← MAX(v, MIN-VALUE(RESULT(s, a)))  
    return v
```

```
function MIN-VALUE(state) returns a utility value  
    if TERMINAL-TEST(state) then return UTILITY(state)  
    v ← ∞  
    for each a in ACTIONS(state) do  
        v ← MIN(v, MAX-VALUE(RESULT(s, a)))  
    return v
```

# The minimax algorithm

- A complete depth-first exploration of the game tree
- Completeness
  - Yes (if tree is finite)
- Optimality
  - Yes (against an optimal opponent)
- Time complexity
  - $O(b^m)$
- Space complexity
  - $O(bm)$  (depth-first exploration)


# The minimax algorithm




# Problem with minimax search

- The number of game states is exponential in the tree's depth
  - Do not examine every node
- Alpha-beta pruning: Prune away branches that cannot possibly influence the final decision
- Bounded lookahead
  - Limit depth for each search
  - This is what chess players do: look ahead for a few moves and see what looks best

# Alpha-beta pruning



# Alpha-beta pruning


$$\begin{aligned}\text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\ &= 3.\end{aligned}$$

# Alpha-beta pruning

**function** ALPHA-BETA-SEARCH(state) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

**return** the action in ACTIONS(state) with value  $v$

**function** MAX-VALUE(state, $\alpha$ , $\beta$ ) **returns** a utility value

**if** TERMINAL-TEST(state) **then return** UTILITY(state)

$v \leftarrow -\infty$

**for each**  $a$  **in** ACTIONS(state) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

**return**  $v$

**function** MIN-VALUE(state, $\alpha$ , $\beta$ ) **returns** a utility value

**if** TERMINAL-TEST(state) **then return** UTILITY(state)

$v \leftarrow +\infty$

**for each**  $a$  **in** ACTIONS(state) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

**if**  $v \leq \alpha$  **then return**  $v$


$\beta \leftarrow \text{MIN}(\beta, v)$

**return**  $v$

# Alpha-beta pruning

- Pruning does not affect the result
- Good move ordering improves effectiveness of pruning
- Killer move heuristic
- Transposition table avoids re-evaluation a state

# Alpha-beta pruning



# Heuristic minimax

- Both minimax and alpha-beta pruning search all the way to terminal states.
  - This depth is usually impractical because moves must be made in a reasonable amount of time (~ minutes).
- Cut off the search earlier with some depth limit
- Use an evaluation function

H-MINIMAX( $s, d$ ) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d+1) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d+1) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

# Evaluation Functions

- The evaluation function should order the terminal states in the same way as the true utility function does
  - States that are wins must evaluate better than draws, which in turn must be better than losses.
- The computation must not take too long!
- For nonterminal states, their orders should be strongly correlated with the actual chances of winning.

# Cutting off search

- Minimax Cutoff is identical to Minimax Value except
  - *Terminal?* is replaced by *Cutoff?*
  - *Utility* is replaced by *Eval*

**if CUTOFF-TEST(state, depth) then return EVAL(state)**

# Stochastic Games

- Uncertain outcomes controlled by chance, not an adversary!
- Why wouldn't we know what the result of an action will be?
  - Explicit randomness: rolling dice
  - Unpredictable opponents: the ghosts respond randomly
  - Actions can fail: when moving a robot, wheels might slip

# Expectimax search



- Values reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- Expectimax search: compute the average score under optimal play
  - Max nodes as in minimax search
  - Chance nodes are like min nodes, but the outcome is uncertain
  - Calculate expected utilities, i.e. take weighted average of children
- The underlying uncertain-result problems can be formulated as Markov Decision Processes

# Expectimax search

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p * value(successor)
    return v
```



$$v = (1/2)(8) + (1/3)(24) + (1/6)(-12) = 10$$

# Expectimax search

- It is possible to perform pruning in expectimax search.
- Common techniques for pruning in expectimax include:
  - **Depth Limiting:** Limiting the depth of the search tree can effectively prune branches that are too deep to be practically explored.
  - **Evaluation Function:** Using an evaluation function to estimate the value of a state without fully exploring its subtree. If the evaluation function indicates that further exploration is unlikely to yield significant improvements, you can prune the subtree.
  - **Probabilistic Pruning:** In scenarios where probabilities are involved, you might prune branches that have very low probabilities of occurring, as they contribute little to the overall expectation.
  - **Iterative Deepening:** Iterative deepening can be combined with pruning techniques to explore deeper parts of the tree only when necessary, based on the current state of the search.

# References

- Stuart Russell and Peter Norvig. 2009. Artificial Intelligence: A Modern Approach (3rd ed.). Prentice Hall Press, Upper Saddle River, NJ, USA.
- Lê Hoài Bắc, Tô Hoài Việt. 2014. Giáo trình Cơ sở Trí tuệ nhân tạo. Khoa Công nghệ Thông tin. Trường ĐH Khoa học Tự nhiên, ĐHQG-HCM.
- Nguyễn Ngọc Thảo, Nguyễn Hải Minh. 2020. Bài giảng Cơ sở Trí tuệ Nhân tạo. Khoa Công nghệ Thông tin. Trường ĐH Khoa học Tự nhiên, ĐHQG-HCM.