

Introduction to Artificial Intelligence

Lecture: Constraint Satisfaction Problems

Outline

- Constraint Satisfaction Problems (CSPs)
- Constraint Propagation: Inference in CSPs
- Backtracking Search for CSPs
- Local Search for CSPs

Constraint Satisfaction Problems (CSPs)

- A constraint satisfaction problem (CSP) uses a factored representation for each state.
 - State = a set of variables and each of which has a value
 - Solution = each variable has a value that satisfies all constraints on that variable
- Take advantage of the structure of states
- General-purpose rather than problem-specific heuristics
 - Identify combinations of variable-value that violate the constraints → eliminate large portions of the search space all at once
 - Solutions to complex problems

Constraint Satisfaction Problems (CSPs)

- A CSP consists of the following three components
 - $X = \{X_1, \dots, X_n\}$: a set of variables
 - $D = \{D_1, \dots, D_n\}$: a set of domains, one for each variable.
 - C : a set of constraints that state allowable combinations of values.
- Each C consists of a pair $\langle \text{scope}, \text{rel} \rangle$
 - scope : a tuple of variables that participate in the constraint
 - A relation rel defines the values that participated variables can take

Constraints in CSPs


- Assume that both X_1 and X_2 have the domain $\{A, B\}$
- “Two variables must have different values”
- A relation can be an explicit list of all tuples of values that satisfy the constraint.
 - $\langle(X_1, X_2), \{(A, B), (B, A)\}\rangle$
- It can be an abstract relation that supports two operations
 - Test whether a tuple is a member of the relation
 - Enumerate the members of the relation
 - $\langle(X_1, X_2), X_1 \neq X_2\rangle$

Solutions for CSPs

- Each state is defined by an assignment of values to some or all the variables.
- A solution to a CSP is a **consistent – complete** assignment.
 - A **consistent** assignment does not violate any constraints.
 - A **complete** assignment has every variable assigned, while a **partial** assignment assigns values to only some variables.

Example: Map Coloring

- Each state is assigned a color in {red, green, blue}.
- Adjacent states have different colors.
- Constraint graph:
 - Nodes \Leftrightarrow Variables
 - Arcs \Leftrightarrow Constraints



Example: Map Coloring

- Variables: $X = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains: $D_i = \{\text{red, green, blue}\}$
- Constraints: Adjacent regions must have different colors

$\{SA \neq WA, SA \neq NT, SA \neq Q, SA \neq NSW,$

$SA \neq V, WA \neq NT, NT \neq Q, Q \neq NSW, NSW \neq V\}$

- $SA \neq WA$ is a shortcut of $\langle(SA, WA), SA \neq WA\rangle$
- $SA \neq WA$ can be fully enumerated as

$\{(\text{red,green}), (\text{red,blue}), (\text{green,red}),$


$(\text{green,blue}), (\text{blue,red}), (\text{blue,green})\}$

Example: Map Coloring

- There are many possible solutions to this problem.

$\{WA = \text{red}, NT = \text{green}, Q = \text{red}, NSW = \text{green},$

$V = \text{red}, SA = \text{blue}, T = \text{red}\}$



Benefits of CSPs

- Provide natural representation for a wide variety of problems
- Many problems intractable in regular state-space search can be solved quickly with CSP formulation.
- Better insights to the problem and its solution.

Variations on the CSP formalism

- Discrete and finite variables
 - n variables, domain size $d \rightarrow O(d^n)$ complete assignments
 - E.g., map coloring, scheduling with time limits, 8-queens etc.
- Discrete, infinite domains
 - Sets of Integers, strings, etc.
 - E.g., job scheduling without deadlines
 - Constraint language: understand constraints without enumeration
- Continuous domains
 - Real-world problems often involve continuous domains and even real-valued variables.

CPSs in practise

- Operations Research (scheduling, timetabling)
 - Scheduling the time of observations on the Hubble Space Telescope
 - Airline schedules
- Linear programming
 - Constraints must be linear equalities or inequalities
→ solved in time polynomial in the number of variables.
- Bioinformatics (DNA sequencing)
- Electrical engineering (circuit layout-ing)
- Cryptography
- Computer vision: image interpretation

Types of constraints

- Unary constraint: restrict the value of a single variable
 - $SA \neq green$
- Binary constraint: relate two variables
 - $SA \neq WA$
- Higher-order constraints: involve three or more variables
 - E.g., Professors A, B, and C cannot be on a committee together
 - Always possible to be represented by multiple binary constraints
- Global constraints: involving an arbitrary number of variables
 - $AllDiff$ = all variables involved must have different values

Preference constraints

- Which solutions are preferred → soft constraints
 - E.g., *red* is better than *green*
→ this often can be represented by a cost for each variable assignment
- Constraint optimization problem (COP): a combination of optimization with CSPs → linear programming

Constraint propagation

- Constraints help to reduce the number of legal values for a variable
 - legal values for another variable are also reduced.
- Intertwined with search, or done as a preprocessing step.
- Enforcing local consistency in each part of a graph causes inconsistent values to be eliminated throughout the graph.



Node consistency

- A single variable is node-consistent if all the values in the variable's domain satisfy the variable's unary constraints.
- Eliminate all the unary constraints in a CSP by running node consistency.
- E.g., The South Australians dislike green, the domain of $\{SA\}$ will be $\{red, \cancel{green}, blue\}$

Arc consistency

- A variable in a CSP is arc-consistent if every value in its domain satisfies the variable's binary constraints.
- Arc consistency may have no effect in several cases.
- E.g., the Australia map, no matter what value chosen for *SA* (or for *WA*), there is a valid value for the other variable.

Arc consistency



Arc consistency

- Run as a preprocessor before the search starts or after each assignment
- AC must be run repeatedly until no inconsistency remains
- Need a systematic method for arc-checking
 - If X loses a value, neighbors of X need to be rechecked
 - incoming arcs can become inconsistent again while outgoing arcs stay still

AC-3 algorithm

function AC-3(csp)

returns false if an inconsistency is found and true otherwise

inputs: csp, a binary CSP with components (X, D, C)

local variables: queue, a queue of arcs, initially all the arcs in csp

while queue is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$

if REVISE(csp, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k in $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to queue

return true

AC-3 algorithm

function REVISE(csp, X_i , X_j)

returns true iff we revise the domain of X_i

revised \leftarrow false

for each x in D_i **do**

if no value y in D_j allows (x ,y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

 revised \leftarrow true

return revised

- The worst-case complexity is $O(cd^3)$
 - n : number of variables, each has domain size d , c binary constraints (arc)

Backtracking Search


- States are defined by the values assigned so far
 - Initial state: empty assignment { }
 - Successor function: assign a value to an unassigned variable that agrees with the current assignment
→ fail if no legal assignments
 - Goal test: the current assignment is complete
- Depth-first search: choose values for one variable at a time and backtrack when a variable has no legal values left

Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return BACKTRACK({ }, csp)
```


```
function BACKTRACK(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
  var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment then
      add {var = value} to assignment
      inferences  $\leftarrow$  INFERENCE(csp, var, value)
      if inferences  $\neq$  failure then
        add inferences to assignment
        result  $\leftarrow$  BACKTRACK(assignment, csp)
        if result  $\neq$  failure then
          return result
    remove {var = value} and inferences from assignment
  return failure
```

Backtracking Search



Variable and value ordering

- Minimum-remaining-values (MRV) heuristic: choose the variable with the fewest legal values
 - E.g., after $[WA = red, NT = green]$ only one possible value for SA




Variable and value ordering

- Degree heuristic (DH): choose the variable that involves in the largest number of constraints on other unassigned variables.
 - E.g., *SA* has a highest degree of 5



Variable and value ordering

- Least constraining value (LCV) heuristic: given a variable, choose the value that leaves the maximum flexibility for subsequent variable assignments



Inference: Forward checking

- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Local search for CSPs

- Complete-state formulation
 - The initial state assigns a value to every variable → violation
 - The search changes the value of one variable at a time → eliminate the violated constraints
- Min-conflicts heuristic: the minimum number of conflicts with other variables
- Min-conflicts is surprisingly effective for many CSPs.

MIN-CONFLICTS algorithm

function MIN-CONFLICTS(*csp*, *max steps*) **returns** a solution or failure

inputs: *csp*, a constraint satisfaction problem

max steps, the number of steps allowed before giving up

current \leftarrow an initial complete assignment for *csp*

for *i* = 1 **to** *max steps* **do**

if *current* is a solution for *csp* **then return** *current*

var \leftarrow a randomly chosen conflicted variable from *csp.VARIABLES*

value \leftarrow the value *v* for *var* that minimizes CONFLICTS(*var*, *v*, *current*, *csp*)

 set *var* = *value* in *current*

return failure

Local search for CSPs

- The landscape of a CSP under the min-conflicts heuristic usually has a series of plateaux.
- Plateau search: allow sideways moves to another state with the same score.
- Tabu search: keep a small list of recently visited states and forbid the algorithm to return to those states
- Simulated annealing can also be used.

Constraint weighting

- Concentrate the search on the important constraints
- Each constraint is given a numeric weight, W , initially all 1.
- At each step, choose a variable/value pair to change that has the lowest total weight of all violated constraints
- Increase the weight of each constraint that is violated by the current assignment.

Homework

- Conduct homework in the given notebook

References

- Stuart Russell and Peter Norvig. 2009. Artificial Intelligence: A Modern Approach (3rd ed.). Prentice Hall Press, Upper Saddle River, NJ, USA.
- Lê Hoài Bắc, Tô Hoài Việt. 2014. Giáo trình Cơ sở Trí tuệ nhân tạo. Khoa Công nghệ Thông tin. Trường ĐH Khoa học Tự nhiên, ĐHQG-HCM.
- Nguyễn Ngọc Thảo, Nguyễn Hải Minh. 2020. Bài giảng Cơ sở Trí tuệ Nhân tạo. Khoa Công nghệ Thông tin. Trường ĐH Khoa học Tự nhiên, ĐHQG-HCM.